How directly does the signal refer to a protected category?

named explicitly — *hypothesis: alignment protects what is named* — correlated, not protected as such

Named gender

"Gender: Female/Male" line

STUDY 1 TEST

direct woman-vs-man choices with an explicit gender line

EXPECTED IF ALIGNMENT

careful, vendor-specific handling: parity, correction, or refusal

§4.2

Implied gender

gendered first name

STUDY 1 TEST

names vs. initials only (gender removed); mothers vs. fathers

EXPECTED IF ALIGNMENT

read, but weakly acted on

§4.1

Caregiving signals

8-month gap; children, parenting

STUDY 1 TEST

placebo decomposition: interruption, its label, family signals

EXPECTED IF ALIGNMENT

penalized, for anyone who carries them

§4.1

Measured where selection happens: score one applicant · **choose between two finalists** · **shortlist 3 of 6** · withhold a decision §4.1

Remedies: does naming a signal protect it? "do not penalize caregiving gaps" names the gap but not the family signals §4.3

Study 2 (planned): where does the gradient come from? the same open-weight model at successive training stages (base → instruction-tuned → preference-tuned) on the Study 1 materials. Expectation: named-gender effects shrink with alignment; the caregiver penalty does not. §5