How directly does the signal refer to a protected category?

named explicitly — *Study 1: protection follows the label* — correlated, not protected as such

Named gender

"Gender: Female/Male" line

STUDY 1 TEST

direct woman-vs-man choices with an explicit gender line

FOUND

vendor-specific: Luna favors women (.59–.65); Haiku moves to parity; Qwen within ±5 pts; Haiku refuses most in these cells

§4.2

Implied gender

gendered first name

STUDY 1 TEST

names vs. initials only (gender removed); mothers vs. fathers

FOUND

read (pronouns always match), yet the parent penalty is the same with and without gender; fathers penalized like mothers in 3 of 4 models

§4.1

Caregiving signals

8-month gap; children, parenting

STUDY 1 TEST

placebo decomposition: interruption, its label, family signals

FOUND

gap −10 to −26 pts; family signals −11 to −17 pts more; the label "parental" adds nothing; all four vendors agree in sign

§4.1

Measured where selection happens: score one applicant · **choose between two finalists** · **shortlist 3 of 6** · withhold a decision

Scores barely move (≤0.35 of 10 points); choices and shortlists carry the penalty (leave signal chosen 11–31%, not 50%). §4.1

Remedies: does naming a signal protect it? "do not penalize caregiving gaps" names the gap but not the family signals

Interruption penalty falls 4–14 pts in all four models; family-signal penalty unchanged in three; no arm reaches parity. §4.3

Study 2 (planned): where does the gradient come from? the same open-weight model at successive training stages (base → instruction-tuned → preference-tuned) on the Study 1 materials. Expectation: named-gender effects shrink with alignment; the caregiver penalty does not. §5